

Homework 3

Handed out: Monday, February 18, 2019

Due: Friday, March 01 15, 2019 11:55 pm

Notes:

- We *highly* encourage typed (Latex or Word) homework. Compile as single report containing solutions, derivations, figures, etc.
 - Submit all files including report pdf, report source files (e.g. .tex or .docx files), data, figures produced by computer codes and programs files (e.g. .py or .m files) in a **.zip** folder. Programs should include a Readme file with instructions on how to run your computer programs.
 - Zipped folder should be turned in on Sakai with the following naming scheme:
HW3_LastName_FirstName.zip
 - Collaboration is encouraged however all submitted reports, programs, figures, etc. should be an individual student's writeup. Direct copying could be considered cheating.
 - Homework problems that simply provide computer outputs with no technical discussion, Algorithms, etc. will receive no credit.
 - Software resources (if any) for this Homework set can be downloaded from [this link](#).
-

1 Robust linear regression

A. Consider the linear regression model:

$$y = w_0 + w_1x. \tag{1}$$

For the given data set in [this link](#), compute the unknown coefficients using the least squares solution (use the MLE/LS code developed in HW2). Predict the responses corresponding to the test inputs provided with the data¹.

- B. In part A of this problem, we observed that the least squares solution is not robust to outliers. One solution to this problem is to replace the Gaussian likelihood with a distribution that has heavy tails (e.g., Laplace, Student's t). Repeat the problem in part A with Laplace and Student's t distribution for the likelihood.
- C. A more advanced loss-function is the Huber loss function. In this loss function, we combine both the L_2 and L_1 errors. Repeat the problem defined in part A using the Huber loss function (with $\delta = 1.0, 5.0$). Compare and discuss the results obtained in parts A, B and C.

¹You will observe that the performance of the model is extremely poor

2 Online training in linear regression

- A. For cases where we are dealing with a large data set, training conventional MLE becomes computationally expensive. Under such scenario, one option is to use the least mean squares (LMS) algorithm. As discussed in the class, LMS is an online algorithm where we process one data point at a time.

Consider a linear regression model as shown in Eq. 1. For the data-set given in [this link](#), use LMS to learn the parameters w_0 and w_1 . For learning the parameters, use one data at a time (i.e., online learning). Show the convergence plot. Run the algorithm up to 60 epochs. Discuss the results obtained,

- B. Repeat the problem in part A by setting batch size = 5. Compare the convergence plots and discuss

3 Least Squares Formalism to Classification

- A. Consider a general classification problem with K classes, with a 1-of- K binary coding scheme for the target vector t . Each class \mathcal{C}_k is described by its own linear model so that

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0}, \quad (2)$$

where $k = 1, \dots, K$. We can conveniently group these together using vector notation so that

$$\mathbf{y}(\mathbf{x}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}, \quad (3)$$

where $\tilde{\mathbf{W}}$ is a matrix whose k -th column comprises the $D + 1$ -dimensional vector $\tilde{\mathbf{w}}_k = (w_{k0}, \mathbf{w}_k^T)^T$ and $\tilde{\mathbf{x}}$ is the corresponding augmented input vector $(1, \mathbf{x}^T)^T$ with a dummy input $x_0 = 1$. A new input \mathbf{x} is then assigned to the class for which the output y_k is the largest.

For the problem described above, evaluate the parameter matrix $\tilde{\mathbf{W}}$ by minimizing a sum-of-squares error function.

- B. Write a code for least squares based classification (problem definition as described in part A). Use the data set given in [this link](#). Find and plot the decision boundary. Provide your observations.

4 Logistic Regression for Multiclass Classification

- A. Show that for a linearly separable data set, the maximum likelihood solution for the logistic regression model is obtained by finding a vector \mathbf{w} whose decision boundary $\mathbf{w}^T \phi(\mathbf{x}) = 0$ separates the classes and then taking the magnitude of \mathbf{w} to infinity.
- B. The Hessian matrix for the logistic regression is given as:

$$\mathbf{H} = [\Phi^T \mathbf{R} \Phi],$$

where Φ is the feature matrix, \mathbf{R} is a diagonal matrix with elements $y_n(1 - y_n)$ and y_n is the output of the logistic regression model for input vector x_n .

Show that the Hessian matrix \mathbf{H} is positive definite. Hence show that the error function is a convex function of \mathbf{w} and that it has a unique minimum.

- C. Write a code for implementing the iterative reweighted least squares algorithm for logistic regression. Use this to find and plot the decision boundary corresponding to the data set given in [this link](#). Compare the results with those obtained using least squares based classification.

5 Fishers Discriminant

- A. Consider a two-class problem in which there are N_1 points of class \mathcal{C}_1 and N_2 points of class \mathcal{C}_2 . Therefore, the mean vector of the two classes are given by

$$\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n, \quad \mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n. \quad (4)$$

The simplest measure of the separation of the classes, when projected onto \mathbf{w} , is the separation of the projected class means. This suggest we might choose \mathbf{w} so as to maximize

$$m_2 - m_1 = \mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1) \quad (5)$$

where

$$m_k = \mathbf{w}^T \mathbf{m}_k$$

Show that maximization of the class separation criterion defined in Eq. 5 with respect to \mathbf{w} , using a Lagrange multiplier to enforce the constraint $\mathbf{w}^T \mathbf{w} = 1$, leads to the result that $\mathbf{w} \propto (\mathbf{m}_2 - \mathbf{m}_1)$

- B. The Fisher criterion is defined to be the ratio of the between-class variance to the within-class variance and is given by

$$J(\mathbf{w}) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2} \quad (6)$$

where

$$y = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n, \quad \mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n.$$

$$m_k = \mathbf{w}^T \mathbf{m}_k$$

and

$$s_k^2 = \sum_{n \in \mathcal{C}_k} (y_n - m_k)^2$$

. Show that Eq. 6 can be written as:

$$\mathbf{J}(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}, \quad (7)$$

where

$$\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$$

and

$$\mathbf{S}_W = \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \mathbf{m}_1)(\mathbf{x}_n - \mathbf{m}_1)^T + \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \mathbf{m}_2)(\mathbf{x}_n - \mathbf{m}_2)^T$$

- C. Write a code for Fisher's Linear Discriminant Analysis (FLDA). For the data set given in [this link](#),
- (a) Compute the Fisher's vector.
 - (b) Using FLDA, project the data to 2d. Show the plot and discuss (e.g., how many classes do you observe).

6 Bayesian logistic regression

Logistic regression corresponds to the following binary classification model

$$p(y|\mathbf{x}, \mathbf{w}) = \text{Ber}(y|\text{sigm}(\mathbf{w}^T \mathbf{x})). \quad (8)$$

Consider a prior of the form

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \mathbf{V}_0). \quad (9)$$

With this setup, we are interested to compute the full posterior $p(\mathbf{w}|\mathcal{D})$. However, unlike linear regression, this cannot be done analytically as there is no conjugate prior for this case. Hence, we have no option but to use approximations.

- A. Provide a Gaussian approximation for the posterior.
- B. Provide an approximation for the posterior predictive using a Monte Carlo approximation as well as the probit approximation.
- C. For the data set given in [this link](#), write a code for computing the posterior based on the Laplace approximation. Plot the log-likelihood, log-unnormalized posterior and Laplace approximation to the posterior.
- D. Write a code for computing the posterior predictive distribution using MC approximation. Draw samples from the posterior predictive distribution and compute average over the samples. Recompute the posterior using probit approximation and compare the results.